

# Assignment 1: KNN Classification from Scratch

**Date:** January 22, 2026

## 1. Introduction

This report details the implementation of a K-Nearest Neighbors (KNN) classifier from scratch to solve two distinct problems: binary classification of breast mass digitized images and multi-class classification using the CIFAR-10 image dataset. The project explores the impact of various distance metrics and different values of  $K$  on model accuracy.

## 2. Task 1: Binary Classification (Breast Cancer Diagnosis)

### 2.1 Methodology

A KNN classifier was implemented using a custom `UniformWeightedKNN` class. The dataset was split into **80% training** and **20% testing** sets. Features were normalized using **Min-Max Scaling** to ensure that distance calculations were not biased by different feature scales.

### 2.2 Mathematical Definitions of Distance Metrics

Let  $\mathbf{a}$  and  $\mathbf{b}$  be two feature vectors of dimension  $n$ :

- **Euclidean Distance ( $L_2$ ):**

$$d(\mathbf{a}, \mathbf{b}) = \sqrt{\sum_{i=1}^n (a_i - b_i)^2}$$

- **Manhattan Distance ( $L_1$ ):**

$$d(\mathbf{a}, \mathbf{b}) = \sum_{i=1}^n |a_i - b_i|$$

- **Minkowski Distance ( $p = 3$ ):**

$$d(\mathbf{a}, \mathbf{b}) = \left( \sum_{i=1}^n |a_i - b_i|^p \right)^{1/p}$$

- **Cosine Similarity (Distance):**

$$d(\mathbf{a}, \mathbf{b}) = 1 - \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\| \|\mathbf{b}\|}$$

- **Hamming Distance:**

$$d(\mathbf{a}, \mathbf{b}) = \frac{1}{n} \sum_{i=1}^n \mathbb{I}(a_i \neq b_i)$$

## 2.3 Experimental Results


The highest accuracy was achieved using **Manhattan** and **Euclidean** metrics at lower  $K$  values.

Metric	K=3	K=4	K=9	K=20	K=47
Euclidean/Minkowski	0.9649	0.9561	0.9561	0.9561	0.9561
Manhattan	0.9649	0.9649	0.9649	0.9561	0.9561
Cosine	0.8860	0.8860	0.8947	0.9123	0.8860
Hamming	0.6930	0.6842	0.6754	0.6754	0.6228


### Best Configuration:

- $K = 3$
- Distance Metric: **Euclidean**
- Test Accuracy: **0.9649**
- Precision: **0.9756** | Recall: **0.9302**

## 2.4 Visualizations



**Figure:** K vs Accuracy for Task 1



**Figure:** Decision Boundary (K=3, Euclidean)

## 2.5 Inferences

- **Metric Sensitivity:**

Euclidean and Manhattan distances performed exceptionally well after normalization. Hamming distance performed poorly because the dataset consists of continuous variables, and binarization leads to significant information loss.

- **Impact of K:**

Smaller  $K$  values (3, 4) captured local structure better. As  $K$  increased to 47, accuracy slightly decreased, indicating a move toward underfitting.

- **Decision Boundary:**

The visualization shows a clear separation between Benign (B) and Malignant (M) classes, though some overlap near the boundary suggests the need for non-linear separation.

### 3. Task 2: Multi-class Classification (CIFAR-10)

#### 3.1 Methodology

The CIFAR-10 dataset was subsampled to **5,000 training** and **2,000 testing** images to optimize computation. Images were normalized to the range **[0, 1]**.


#### 3.2 Experimental Results

Metric	K=1	K=3	K=5	K=7	K=9
Manhattan	0.2835	0.2740	0.2950	<b>0.3085</b>	0.2995
Euclidean	0.2630	0.2625	0.2685	0.2845	0.2770
Cosine	0.2855	0.2605	0.2755	0.2870	0.2870
Minkowski	0.2335	0.2365	0.2555	0.2565	0.2525
Hamming	0.2180	0.2070	0.2340	0.2395	0.2430

##### Best Configuration:

- $K = 7$
- Distance Metric: **Manhattan**
- Test Accuracy: **0.3085**
- Precision (Macro): **0.3801** | Recall (Macro): **0.3087**

### 3.3 Visualizations



**Figure:** K vs Accuracy for CIFAR-10

### 3.4 Inferences

- **Metric Superiority:**

Manhattan distance performed better than Euclidean for image data, which is common in high-dimensional spaces.

- **Data Complexity:**

The lower accuracy (~30.8%) highlights that raw pixel values are not very discriminative for complex image classification without feature extraction techniques like CNNs.

- **Optimal K:**

$K = 7$  provided the best balance between noise reduction and class discrimination.

## 4. Conclusion

This assignment demonstrates that KNN is a strong baseline model but highly dependent on distance metrics and feature scaling. While it performs well on structured tabular data (Task 1), its limitations become clear on high-dimensional image data (Task 2), emphasizing the need for more advanced models like CNNs for image classification.