

Group_E_Assignment02_Report

Task 1: Optimizer Performance on Non-Convex Functions

Objective

The objective of this task is to study the behavior of different optimization algorithms on non-convex functions and analyze their convergence characteristics, sensitivity to learning rate, and final solution quality.

Problem Setup

We considered two non-convex functions:

1. Rosenbrock function: $f(x, y) = (1-x)^2 + 100(y - x^2)^2$
2. Discontinuous function: $f(x) = \sin(1/x)$, with $f(0)=0$

The following optimizers were implemented from scratch in Python:

- Gradient Descent (GD)
- Stochastic Gradient Descent with Momentum
- RMSProp
- Adagrad
- Adam

Learning rates tested: 0.01, 0.05, and 0.1.



Stopping criterion: optimization terminated when the change in function value between successive iterations fell below a predefined threshold.

Methodology


For each optimizer and learning rate:

- Gradients were computed analytically.
- Parameters were updated iteratively.
- Convergence behavior (loss vs. iterations) was recorded.
- Final optimized values of x (and y where applicable) were reported.
- Time taken for convergence was measured.

Results and Analysis



Adam: Effect of Learning Rate (Rosenbrock)



- Compare speed of convergence across optimizers.

Relative convergence speed (from fastest to slowest, in general):

1. Adam — Fastest overall

- Typically shows the **quickest drop in loss in early iterations**.
- Handles non-convex landscapes well due to adaptive moment estimates.

2. RMSProp — Very fast

- Usually converges almost as quickly as Adam.
- More stable than plain SGD on rugged loss surfaces.

3. SGD with Momentum — Moderate speed

- Faster than vanilla Gradient Descent because it builds velocity and helps escape shallow regions and local minima.
- Performs well but still sensitive to learning rate choice.

4. Adagrad — Fast initially, slow later

- Often reduces loss quickly in early iterations.
- Becomes very slow later due to continuously shrinking effective learning rate.

5. Vanilla Gradient Descent — Slowest

- Requires the most iterations to converge.
- Highly sensitive to learning rate and often struggles with non-convex functions like Rosenbrock.

- Discuss sensitivity to learning rate.

Impact of learning rate on convergence:

- $\alpha = 0.01$ (**small learning rate**):
 - Generally **stable but slow convergence**.
 - Loss decreases smoothly but requires many iterations to reach the optimum.
 - Less risk of divergence, even for difficult functions like Rosenbrock.
- $\alpha = 0.05$ (**moderate learning rate**):
 - Often provides the **best balance between speed and stability**.
 - Typically leads to faster convergence without major oscillations.
- $\alpha = 0.1$ (**large learning rate**):
 - Can cause **faster initial progress** but risks **oscillations or divergence**, especially for Gradient Descent and Momentum-based SGD.
 - More likely to overshoot the minimum in highly curved regions of the loss surface.

Optimizer robustness to learning rate:

- Adam and RMSProp are **less sensitive** to learning rate changes due to adaptive scaling of updates.
- Vanilla GD and Momentum SGD are **highly sensitive**, requiring careful tuning of α .
- Overall, the results highlight that choosing an appropriate learning rate is crucial, especially for non-convex functions with multiple local minima.

Task 2: Linear Regression Using Multi-Layer Neural Network (From Scratch)

Dataset

Boston Housing Dataset was used with two features:

- RM (average number of rooms per dwelling)
 - CRIM (per capita crime rate by town)
- Target variable: MEDV (median value of homes).

Preprocessing

- Features were normalized.
- Data was split into 80% training and 20% testing.

Model Architecture

- Input Layer: 2 neurons
- Hidden Layer 1: 5 neurons (ReLU activation)

- Hidden Layer 2: 3 neurons (ReLU activation)
- Output Layer: 1 neuron (linear activation)

Training Setup

Optimizers implemented:

- Basic Gradient Descent
- Momentum
- Adam


Learning rates tested: 0.01 and 0.001

Epochs: 1000


Loss function: Mean Squared Error (MSE)

Evaluation

- Test MSE reported
Test Mean Squared Error: 24.40887101467112
- Plot: Predicted vs Actual values



- Loss vs Epoch plot



Bonus Experiments

- Added a third hidden layer with 2 neurons and analyzed impact.
- Implemented L2 regularization and discussed effect on overfitting.

Task 3: Multi-class Classification Using FCNN

Dataset (Custom)

Two datasets were used:

1. Linearly separable dataset (3 classes, 2D)
2. Non-linearly separable dataset (2 or 3 classes, 2D)


Each class was split into:

- 60% Training
- 20% Validation
- 20% Testing


Model

- Dataset 1: FCNN with 1 hidden layer
- Dataset 2: FCNN with 2 hidden layers
- Loss: Squared error
- Training: Stochastic Gradient Descent (SGD)


Results (Best Architecture)




Nonlinear Dataset Decision Region



Training Layer 1 Node 0



Validation Layer 1 Node 0



- Validation Confusion matrix:

```
[[ 1  0 99]
 [44  0 56]
 [ 2  0 98]]
```

- Validation Accuracy: 0.33

- Test Confusion matrix:

```
[[ 0  0 100]
 [ 47  0 53]
 [  8  0 92]]
```

- Test Accuracy: 0.30666666666666664

Task 4: FCNN on MNIST (Optimizer Comparison)

Dataset (MNIST)

Subset of MNIST with 5 chosen classes.

- Image size: 28×28 → flattened to 784-dimensional vector.

- Split: 80% training, 20% testing.


Model (PyTorch)

FCNN with 3 to 5 hidden layers.


Loss: Cross-entropy.

Results

- The **3-hidden-layer model** converged fastest but slightly underperformed in validation accuracy.



- The **4-hidden-layer model (arch_4h)** achieved the best balance between convergence speed and generalization, making it the optimal choice.



- The **5-hidden-layer model** did not provide significant improvement over arch_4h and in some cases showed signs of overfitting or slower convergence.



BEST RESULT

- Best Architecture: arch_4h
- Hidden Layer Sizes: [512, 256, 128, 64]
- Best Optimizer: SGD_momentum
- Validation Confusion Matrix:

```
[[5496,      0,      0,      1,      1],
 [  0, 6322,      0,      3,      1],
 [  0,      0, 5631,     12,      0],
 [  0,      0,      3, 5656,      0],
 [  0,      1,      6,      1, 5454]]
```

- Validation Accuracy: 0.9935637330348398
- Test Confusion Matrix:

```
[[1399,      0,      3,      2,      1],
 [  0, 1542,      3,      1,      5],
 [  5,      2, 1325,     12,      3],
 [  1,      1,      2, 1478,      0],
 [  2,      1,      1,      1, 1357]]
```

- Training Accuracy: 0.998985588358752

Optimizers Compared

Across all architectures (3h, 4h, and 5h), the loss curves demonstrate distinct convergence characteristics for different optimizers:

- **SGD (batch_size=1):**

Exhibited the slowest convergence among all methods, with noisy and fluctuating loss curves, particularly in deeper architectures. This is expected due to high variance in gradient estimates.

- **Batch Gradient Descent:**

Showed smoother loss curves than SGD but converged more slowly in terms of epochs, especially for the 5-hidden-layer architecture.

- **SGD with Momentum (0.9):**

Provided significantly faster and more stable convergence compared to vanilla SGD. This optimizer performed best overall and was used in the final selected model.

- **NAG Momentum:**

Demonstrated behavior similar to standard momentum but with slightly improved stability in some architectures, particularly in early epochs.

- **RMSProp:**

Converged relatively fast in shallow architectures but showed some instability in deeper networks due to adaptive step sizes.

- **Adam:**

Showed rapid initial convergence but in some cases plateaued earlier than momentum-based methods, leading to slightly lower validation performance compared to SGD with momentum.

Conclusion

This assignment demonstrates the impact of optimizers, network depth, and regularization on convergence and generalization. Task 1 highlights optimizer behavior on non-convex landscapes, Task 2 shows the effectiveness of multi-layer networks for regression, Task 3 explores decision boundaries in classification, and Task 4 provides a comprehensive comparison of optimizers on MNIST.